

# HotSpot

Romanian National AI Olympiad 2025 - Problem Editorial

AI Olympiad Committee

June 23, 2025

## Abstract

This editorial presents a comprehensive analysis of the HotSpot problem, an image segmentation challenge set in a dystopian future where WiFi signals represent areas of clean communication. The problem requires participants to identify geometric shapes (HotSpots) in noisy satellite imagery while dealing with various levels of interference and background noise. We explore multiple solution approaches ranging from simple thresholding to advanced region growing algorithms, analyzing their effectiveness across different subtasks.

## Contents

<b>1</b>	<b>Problem Overview</b>	<b>2</b>
1.1	Context and Motivation . . . . .	2
1.2	Technical Specification . . . . .	2
1.3	Geometric Shapes . . . . .	2
<b>2</b>	<b>Suggested Solution Approaches</b>	<b>2</b>
2.1	Baseline Solutions . . . . .	2
2.1.1	Simple Thresholding (Subtasks 1-2) . . . . .	2
2.1.2	K-Means Clustering . . . . .	3
2.2	Advanced Suggestions . . . . .	3
2.2.1	Region Growing with Similarity-Based Expansion . . . . .	3
2.2.2	Cellular Automata Segmentation . . . . .	4
2.2.3	Brightness Thresholding with Morphological Operations . . . . .	6
2.3	Machine Learning Suggestion . . . . .	6
2.3.1	Training Data Preparation . . . . .	6
2.3.2	Segmentation Classifier . . . . .	7
<b>3</b>	<b>Handling Noise and Interference</b>	<b>7</b>
3.1	Types of Interference . . . . .	7
3.2	Noise Reduction Strategies . . . . .	8
3.2.1	Gaussian Filtering . . . . .	8
3.2.2	Median Filtering . . . . .	8
3.2.3	Stripe Detection and Removal . . . . .	8
<b>4</b>	<b>Suggested Implementation Strategy</b>	<b>8</b>
4.1	Progressive Approach . . . . .	8
4.2	Parameter Tuning Suggestions . . . . .	9
4.3	Validation Strategy . . . . .	9
<b>5</b>	<b>Conclusion</b>	<b>9</b>

# 1 Problem Overview

## 1.1 Context and Motivation

The HotSpot problem is set in the year 2147, where humanity has transitioned to a completely wireless society. In this dystopian world, WiFi signals serve as the primary source of energy, communication, and control. The challenge lies in identifying clean signal areas (HotSpots) from satellite imagery that is corrupted by radio interference and background noise.

We designed this problem to be something unique rather than just applying classical deep learning models. While HotSpot is fundamentally a segmentation problem dealing with basic geometric shapes, the variety of possible solution approaches is staggering. This diversity allows participants to explore different algorithmic paradigms and discover that sometimes simpler, more traditional computer vision techniques can outperform modern deep learning approaches.

## 1.2 Technical Specification

Participants must process four sets of  $256 \times 256$  RGB images:

- **Satellite\_Images-1:** Images with at least one object, minimal noise
- **Satellite\_Images-2:** Images with multiple objects, minimal noise
- **Satellite\_Images-3:** Images with moderate complexity
- **Satellite\_Images-4:** Images with maximum complexity and interference

The goal is to generate binary masks identifying HotSpots, encoded using Run-Length Encoding (RLE) for efficient storage and evaluation.

## 1.3 Geometric Shapes

HotSpots appear as various geometric primitives:

- Circles and ellipses
- Rectangles and squares
- Convex and concave polygons

# 2 Suggested Solution Approaches

## 2.1 Baseline Solutions

### 2.1.1 Simple Thresholding (Subtasks 1-2)

For the first two subtasks, where images contain minimal noise, we suggest a straightforward approach that works effectively:

---

**Algorithm 1** Simple Non-Black Pixel Detection

---

**Require:** RGB image  $I$  of size  $256 \times 256$

**Ensure:** Binary mask  $M$

```
1: for each pixel  $(i, j)$  in  $I$  do
2:   if  $I[i, j] \neq (0, 0, 0)$  then
3:      $M[i, j] \leftarrow 1$ 
4:   else
5:      $M[i, j] \leftarrow 0$ 
6:   end if
7: end for
8: return  $M$ 
```

---

This approach leverages the clean nature of early subtasks where HotSpots are clearly distinguished from the black background.

### 2.1.2 K-Means Clustering

A more sophisticated baseline suggestion uses K-means clustering to separate foreground objects from background:

Listing 1: K-Means Segmentation

```
1 import numpy as np
2 from sklearn.cluster import KMeans
3
4 def kmeans_segmentation(image, k=2):
5     # Reshape image to pixel array
6     pixels = image.reshape(-1, 3)
7
8     # Apply K-means clustering
9     kmeans = KMeans(n_clusters=k, random_state=42)
10    labels = kmeans.fit_predict(pixels)
11
12    # Reshape back to image dimensions
13    segmented = labels.reshape(image.shape[:2])
14
15    # Identify foreground cluster (non-black)
16    cluster_means = kmeans.cluster_centers_
17    foreground_cluster = np.argmax(np.sum(cluster_means, axis=1))
18
19    return (segmented == foreground_cluster).astype(int)
```

## 2.2 Advanced Suggestions

### 2.2.1 Region Growing with Similarity-Based Expansion

One of the most effective suggested approaches for complex subtasks employs region growing based on pixel similarity within local neighborhoods:

---

**Algorithm 2** Region Growing Segmentation

---

**Require:** RGB image  $I$ , similarity threshold  $\tau$ , neighborhood size  $N$

**Ensure:** Binary mask  $M$

```
1: Initialize  $M \leftarrow \mathbf{0}$ 
2: Initialize seed points  $S$  using brightness thresholding
3: for each seed point  $s \in S$  do
4:   Initialize queue  $Q \leftarrow \{s\}$ 
5:    $M[s] \leftarrow 1$ 
6:   while  $Q \neq \emptyset$  do
7:      $p \leftarrow \text{dequeue}(Q)$ 
8:     for each neighbor  $n$  of  $p$  in  $N$ -neighborhood do
9:       if  $M[n] = 0$  AND  $\text{similarity}(I[p], I[n]) > \tau$  then
10:         $M[n] \leftarrow 1$ 
11:         $\text{enqueue}(Q, n)$ 
12:       end if
13:     end for
14:   end while
15: end for
16: return  $M$ 
```

---

The similarity function can be defined as:

$$\text{similarity}(p_1, p_2) = \exp\left(-\frac{\|p_1 - p_2\|_2^2}{2\sigma^2}\right)$$

where  $\sigma$  controls the sensitivity to color differences.

### 2.2.2 Cellular Automata Segmentation

An interesting alternative approach employs cellular automata for image segmentation. This method treats each pixel as a cell in a 2D grid, where the cell's state evolves based on local neighborhood rules. Surprisingly, this bio-inspired approach often produces results very similar to K-means clustering while offering different computational characteristics.

---

**Algorithm 3** Cellular Automata Segmentation

---

**Require:** RGB image  $I$ , iterations  $T$ , threshold  $\theta$

**Ensure:** Binary mask  $M$

```
1: Initialize state matrix  $S$  based on pixel brightness
2: for  $t = 1$  to  $T$  do
3:    $S_{\text{new}} \leftarrow S$ 
4:   for each pixel  $(i, j)$  in  $I$  do
5:     Count active neighbors  $N_{\text{active}} \leftarrow \sum_{(x,y) \in \mathcal{N}(i,j)} S[x, y]$ 
6:     Compute local brightness  $B \leftarrow \|I[i, j]\|_2$ 
7:     if  $B > \theta$  AND  $N_{\text{active}} \geq 2$  then
8:        $S_{\text{new}}[i, j] \leftarrow 1$  {Cell becomes active}
9:     else if  $B \leq \theta$  AND  $N_{\text{active}} < 3$  then
10:       $S_{\text{new}}[i, j] \leftarrow 0$  {Cell becomes inactive}
11:     end if
12:   end for
13:    $S \leftarrow S_{\text{new}}$ 
14: end for
15:  $M \leftarrow S$ 
16: return  $M$ 
```

---

The cellular automata approach uses simple local rules that create emergent segmentation behavior. The update rules can be customized based on the specific characteristics of the HotSpots:

Listing 2: Cellular Automata Implementation

```

1 import numpy as np
2 from scipy.ndimage import convolve
3
4 def cellular_automata_segmentation(image, iterations=10,
5     brightness_threshold=0.3):
6     # Convert to grayscale and normalize
7     gray = np.mean(image, axis=2) / 255.0
8
9     # Initialize cellular automata state
10    state = (gray > brightness_threshold).astype(int)
11
12    # Define neighborhood kernel (Moore neighborhood)
13    kernel = np.array([[1, 1, 1],
14                       [1, 0, 1],
15                       [1, 1, 1]])
16
17    for iteration in range(iterations):
18        # Count active neighbors
19        neighbor_count = convolve(state, kernel, mode='constant', cval
20            =0)
21
22        # Apply cellular automata rules
23        new_state = state.copy()
24
25        # Rule 1: High brightness pixels with sufficient neighbors
26        # become active
27        activate_mask = (gray > brightness_threshold) & (neighbor_count
28            >= 2)
29        new_state[activate_mask] = 1
30
31        # Rule 2: Low brightness pixels with few neighbors become
32        # inactive
33        deactivate_mask = (gray <= brightness_threshold) & (
34            neighbor_count < 3)
35        new_state[deactivate_mask] = 0
36
37        # Rule 3: Stability rule - maintain state if conditions are
38        # ambiguous
39        stable_mask = (neighbor_count >= 3) & (neighbor_count <= 5)
40        new_state[stable_mask] = state[stable_mask]
41
42        state = new_state
43
44    return state

```

Interestingly, the cellular automata approach often converges to segmentation results that are remarkably similar to those obtained through K-means clustering. This similarity arises because both methods attempt to group pixels with similar characteristics, though through fundamentally different mechanisms: K-means through explicit distance minimization in feature space, and cellular automata through local neighborhood interactions that create emergent clustering behavior.

The key advantages of cellular automata segmentation include:

- **Inherent spatial coherence:** The local neighborhood rules naturally preserve spatial relationships
- **Noise robustness:** The iterative nature helps filter out isolated noise pixels
- **Computational efficiency:** Simple local operations can be easily parallelized
- **Adaptability:** Rules can be easily modified for different image characteristics

### 2.2.3 Brightness Thresholding with Morphological Operations

We suggest this approach that combines intensity-based segmentation with morphological post-processing:

Listing 3: Thresholding with Morphology

```

1 import cv2
2 from skimage import morphology
3
4 def threshold_morphology_segmentation(image, threshold=0.3):
5     # Convert to grayscale and normalize
6     gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
7     normalized = gray.astype(float) / 255.0
8
9     # Apply adaptive thresholding
10    binary_mask = normalized > threshold
11
12    # Morphological operations
13    # Remove noise
14    cleaned = morphology.remove_small_objects(binary_mask, min_size=50)
15
16    # Fill holes
17    filled = morphology.remove_small_holes(cleaned, area_threshold=100)
18
19    # Smooth boundaries
20    kernel = morphology.disk(2)
21    final_mask = morphology.closing(filled, kernel)
22
23    return final_mask.astype(int)

```

## 2.3 Machine Learning Suggestion

For subtasks 3 and 4, a supervised learning approach can be considered, though interestingly, this problem demonstrates that modern deep learning architectures like U-Net do not perform well due to the specific nature of the geometric shapes and noise patterns involved.

### 2.3.1 Training Data Preparation

Using the clean results from subtasks 1 and 2 as training data:

Listing 4: Feature Extraction for ML

```

1 def extract_features(image, window_size=5):
2     features = []
3     h, w = image.shape[:2]
4
5     for i in range(window_size//2, h - window_size//2):
6         for j in range(window_size//2, w - window_size//2):
7             # Extract local window

```

```

8         window = image[i-window_size//2:i+window_size//2+1,
9                        j-window_size//2:j+window_size//2+1]
10
11         # Compute features
12         mean_rgb = np.mean(window, axis=(0,1))
13         std_rgb = np.std(window, axis=(0,1))
14         local_contrast = np.std(cv2.cvtColor(window, cv2.
15                                COLOR_RGB2GRAY))
16
17         features.append(np.concatenate([mean_rgb, std_rgb, [
18                                local_contrast]]))
19
20     return np.array(features)

```

### 2.3.2 Segmentation Classifier

A Random Forest or SVM classifier can be trained on these features:

Listing 5: Segmentation Classifier

```

1 from sklearn.ensemble import RandomForestClassifier
2 from sklearn.svm import SVC
3
4 def train_segmentation_classifier(train_images, train_masks):
5     X, y = [], []
6
7     for img, mask in zip(train_images, train_masks):
8         features = extract_features(img)
9         labels = mask.flatten()
10
11         X.extend(features)
12         y.extend(labels)
13
14     # Train classifier
15     classifier = RandomForestClassifier(n_estimators=100, random_state
16                                     =42)
17     classifier.fit(X, y)
18
19     return classifier
20
21 def predict_segmentation(classifier, image):
22     features = extract_features(image)
23     predictions = classifier.predict(features)
24
25     return predictions.reshape(image.shape[:2])

```

## 3 Handling Noise and Interference

### 3.1 Types of Interference

The problem specifies two main sources of interference:

- **Random background noise:** Gaussian or salt-and-pepper noise
- **Semi-transparent stripes:** Linear interference patterns

## 3.2 Noise Reduction Strategies

### 3.2.1 Gaussian Filtering

Apply Gaussian blur to reduce high-frequency noise:

$$G_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

### 3.2.2 Median Filtering

Effective for salt-and-pepper noise removal while preserving edges.

### 3.2.3 Stripe Detection and Removal

For semi-transparent stripes, use Hough transform or directional filtering:

Listing 6: Stripe Removal

```
1 def remove_strips(image, angle_range=(-90, 90)):  
2     # Convert to grayscale  
3     gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)  
4  
5     # Apply Hough line detection  
6     lines = cv2.HoughLines(gray, 1, np.pi/180, threshold=100)  
7  
8     # Create mask for detected lines  
9     line_mask = np.zeros_like(gray)  
10  
11     if lines is not None:  
12         for line in lines:  
13             rho, theta = line[0]  
14             a = np.cos(theta)  
15             b = np.sin(theta)  
16             x0 = a * rho  
17             y0 = b * rho  
18  
19             # Draw line on mask  
20             pt1 = (int(x0 + 1000*(-b)), int(y0 + 1000*(a)))  
21             pt2 = (int(x0 - 1000*(-b)), int(y0 - 1000*(a)))  
22             cv2.line(line_mask, pt1, pt2, 255, 2)  
23  
24     # Inpaint to remove stripes  
25     result = cv2.inpaint(image, line_mask, 3, cv2.INPAINT_TELEA)  
26     return result
```

## 4 Suggested Implementation Strategy

### 4.1 Progressive Approach

We recommend the following progression:

1. **Subtasks 1-2:** Consider using simple thresholding to establish baseline results
2. **Subtask 3:** Try applying region growing or supervised learning trained on 1-2
3. **Subtask 4:** Experiment with combining noise reduction with advanced segmentation



## 4.2 Parameter Tuning Suggestions

Key parameters that may require optimization:

- Similarity threshold  $\tau$  for region growing
- Brightness threshold for initial segmentation
- Morphological kernel sizes
- Noise reduction filter parameters
- Number of iterations for cellular automata

## 4.3 Validation Strategy

Consider using cross-validation on subtasks 1-2 to optimize parameters before applying to harder subtasks.

## 5 Conclusion

The HotSpot problem demonstrates that effective image segmentation doesn't always require the most sophisticated deep learning models. Through progressive difficulty across subtasks, participants can explore various algorithmic approaches and discover that traditional computer vision techniques often provide robust solutions for geometric shape segmentation tasks.